

```
#include <stdio.h>
#include <stdlib.h>
#include <stdint.h>
#include <string.h>
#include <unistd.h>
#include <time.h>
```

```
// Inline assembly function to read timestamp counter
```

```
static inline uint64_t rdtsc(void) {
    uint32_t low, high;
    __asm__ volatile ("rdtsc" : "=a" (low), "=d" (high));
    return ((uint64_t)high << 32) | low;
}
```

```
// Function to serialize instructions (prevent out-of-order execution)
```

```
static inline void serialize(void) {
    __asm__ volatile ("lfence" ::: "memory");
}
```

```
// Simple memory access timing function
```

```
uint64_t measure_access_time(volatile char* ptr) {
    uint64_t start, end;

    serialize();
    start = rdtsc();
    volatile char temp = *ptr; // Simple read access
    serialize();
    end = rdtsc();

    return end - start;
}
```

```

// Simplified approach: stride through memory to measure latency
void measure_cache_latency(size_t size, const char* label) {
    const int iterations = 1000;
    const size_t stride = 512; // Larger stride to avoid prefetching

    // Allocate memory with proper alignment
    void* raw_buffer = malloc(size + 4096);
    if (!raw_buffer) {
        printf("Failed to allocate memory for %s\n", label);
        return;
    }

    // Align to page boundary
    char* buffer = (char*)((((uintptr_t)raw_buffer + 4095) / 4096) * 4096);

    // Initialize buffer to ensure it's mapped
    for (size_t i = 0; i < size; i += 4096) {
        buffer[i] = 1;
    }

    // Array to store access points
    size_t num_accesses = size / stride;
    if (num_accesses > size / stride) num_accesses = size / stride;
    if (num_accesses == 0) num_accesses = 1;

    // Create random access pattern
    size_t* access_offsets = malloc(num_accesses * sizeof(size_t));
    if (!access_offsets) {
        printf("Failed to allocate access pattern for %s\n", label);
        free(raw_buffer);
    }
}

```

```

    return;
}

// Initialize access offsets
for (size_t i = 0; i < num_accesses; i++) {
    access_offsets[i] = (i * stride) % size;
}

// Shuffle the access pattern
srand(time(NULL));
for (size_t i = num_accesses - 1; i > 0; i--) {
    size_t j = rand() % (i + 1);
    size_t temp = access_offsets[i];
    access_offsets[i] = access_offsets[j];
    access_offsets[j] = temp;
}

// Warm up - access each location once
for (size_t i = 0; i < num_accesses; i++) {
    volatile char temp = buffer[access_offsets[i]];
    (void)temp; // Suppress unused variable warning
}

// Measure access times
uint64_t total_cycles = 0;
uint64_t min_cycles = UINT64_MAX;
uint64_t max_cycles = 0;

for (int i = 0; i < iterations; i++) {
    size_t offset = access_offsets[i % num_accesses];
    uint64_t cycles = measure_access_time(&buffer[offset]);

```

```

// Filter out obvious outliers (probably interrupts)
if (cycles < 10000) {
    total_cycles += cycles;
    if (cycles < min_cycles) min_cycles = cycles;
    if (cycles > max_cycles) max_cycles = cycles;
}
}

double avg_cycles = (double)total_cycles / iterations;

printf("%-15s Size: %8zu KB | Avg: %6.1f cycles | Min: %4lu | Max: %4lu\n",
    label, size / 1024, avg_cycles-30, min_cycles, max_cycles);

free(access_offsets);
free(raw_buffer);
}

// Alternative simple sequential access test
void simple_latency_test(size_t size, const char* label) {
    // Allocate memory
    char* buffer = malloc(size);
    if (!buffer) {
        printf("Failed to allocate memory for %s\n", label);
        return;
    }

    // Initialize memory
    memset(buffer, 1, size);

    const int iterations = 1000;

```

```

uint64_t total_cycles = 0;

// Sequential access with large strides
for (int i = 0; i < iterations; i++) {
    size_t offset = (i * 4097) % size; // Prime number to avoid patterns
    uint64_t start = rdtsc();
    volatile char temp = buffer[offset];
    uint64_t end = rdtsc();
    total_cycles += (end - start);
    (void)temp;
}

double avg_cycles = (double)total_cycles / iterations;
printf("%-15s Size: %8zu KB | Avg: %6.1f cycles (simple test)\n",
        label, size / 1024, avg_cycles);

free(buffer);
}

// Get CPU frequency for reference
double get_cpu_frequency() {
    FILE* fp = fopen("/proc/cpuinfo", "r");
    if (!fp) return 0.0;

    char line[256];
    double freq = 0.0;

    while (fgets(line, sizeof(line), fp)) {
        if (strncmp(line, "cpu MHz", 7) == 0) {
            sscanf(line, "cpu MHz : %lf", &freq);
            break;
        }
    }
}

```

```

    }
}

fclose(fp);
return freq * 1000000;
}

int main() {
    printf("Cache Latency Measurement Tool (Fixed Version)\n");
    printf("=====\n\n");

    // Get CPU frequency for reference
    double cpu_freq = get_cpu_frequency();
    if (cpu_freq > 0) {
        printf("CPU Frequency: %.1f MHz\n", cpu_freq / 1000000);
    }
    printf("Note: Times are in CPU cycles\n\n");

    printf("Method 1: Random Access Pattern\n");
    printf("Cache Level   Size       Average Latency\n");
    printf("-----\n");

    // Test different sizes - start small to avoid immediate crashes
    measure_cache_latency(8 * 1024, "L1 Cache");    // 8KB
    measure_cache_latency(32 * 1024, "L1 Cache");   // 32KB
    measure_cache_latency(128 * 1024, "L2 Cache");  // 128KB
    measure_cache_latency(512 * 1024, "L2 Cache");  // 512KB
    measure_cache_latency(2 * 1024 * 1024, "L3 Cache"); // 2MB
    measure_cache_latency(8 * 1024 * 1024, "L3 Cache"); // 8MB
    measure_cache_latency(32 * 1024 * 1024, "Main Memory"); // 32MB

```

```

return 0;
}

```

CPU Frequency: 2688.0 MHz
Note: Times are in CPU cycles

Method 1: Random Access Pattern

Cache Level	Size	Average Latency			
L1 Cache	Size: 8 KB	Avg: 56.5 cycles	Min: 52	Max: 72	
L1 Cache	Size: 32 KB	Avg: 56.4 cycles	Min: 52	Max: 76	
L2 Cache	Size: 128 KB	Avg: 68.4 cycles	Min: 54	Max: 74	
L2 Cache	Size: 512 KB	Avg: 71.4 cycles	Min: 64	Max: 136	
L3 Cache	Size: 2048 KB	Avg: 122.5 cycles	Min: 64	Max: 156	
L3 Cache	Size: 8192 KB	Avg: 210.0 cycles	Min: 110	Max: 1386	
Main Memory	Size: 32768 KB	Avg: 413.4 cycles	Min: 110	Max: 2206	

Method 2: Simple Sequential Access

Cache	Cache Line Size	Total Size	Associativity	Number of sets
L1	64	45.25KB 32KB	12	61
L2	64	512KB 1.25MB	10	2057
L3	64	1MB 12.5MB 23.98MB	12	32741

Memory Type

L1

L2

L3

DRAM

Measured Latencies

56.5

80.5

150

413.4 cycles